

1. (3 points) Bagging can only be applied to decision trees.
  - (a) True
  - (b) False
  - (c) It depends
  - (d) All of the above
  - (e) None of the above
2. (3 points) Gradient Boosting can be applied on regression and classification tasks.
  - (a) True
  - (b) False
  - (c) It depends
  - (d) All of the above
  - (e) None of the above
3. (3 points) Which of the following **is/are** true in the Gradient Boosting model?
  - (a) The decision tree learners are build in sequence.
  - (b) The decision tree learners are build in parallel.
  - (c) The decision tree learners are build in sequence and parallel at the same time.
  - (d) (a) and (b)
  - (e) (a) and (c)
  - (f) (b) and (c)
  - (g) (a), (b) and (c)

Consider the `Customer_Churn.csv` datafile. Each row represents a customer, each column contains customer's attributes described on the column Metadata. The data set includes information about:

- Customers who left within the last month, the column is called Churn.
- Services that each customer has signed up for phone, multiple lines, internet, online security, online backup, device protection, tech support, and streaming TV and movies.
- Customer account information: how long they've been a customer, contract, payment method, paperless billing, monthly charges, and total charges.
- Demographic info about customers: gender, age range, and if they have partners and dependent

4. **In Python**, answer the following:

- (a) (3 points) Using the `pandas` library, read the csv file and create a data-frame called `churn`.
- (b) (3 points) Using the `numpy` library, create a variable called `Churn_numb` that takes the value of 1 when `Churn = Yes` and 0 when `Churn = No`.
- (c) (5 points) Using the `sklearn.ensemble` library, build a gradient boosting classifier, in which `Churn_numb` is the target variable, and `tenure` and `MonthlyCharges` are the input variables. Predict the likelihood of churn of a customer with `tenure` equal to 20, and `MonthlyCharges` equal to 80.